



SPEECH EMOTIONAL RECOGNITION

Assignment-3 Patter Recognition

Ahmed Abdelkabier Hassan	6669
Mohanned Bashar	6656

1-Download Data set and understand it:

```
import os
if not os.path.exists(r'/content/feature_space.pkl'):
    audio_signal = []
    samples_rate = []
    files_path = pd.Series(files_path , name='.wav_path').astype(str)
    labels = pd.Series(emo_labels , name='labels')

    for audios in files_path:
        audio , sample_rate = lib.load(audios)
        audio_signal.append(audio)
        samples_rate.append(sample_rate)


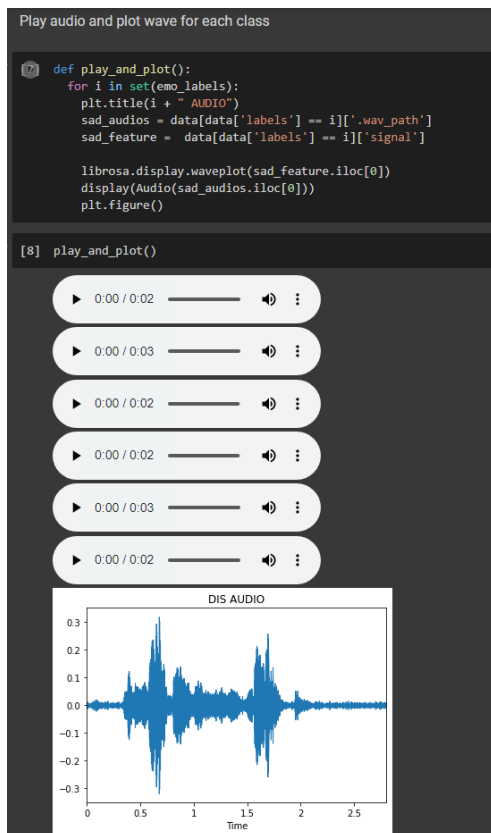
    audio_features = pd.Series(audio_signal , name='signal')
    data = pd.concat([files_path,audio_features,labels] , axis=1)
else:
    data = None
```

[5] data

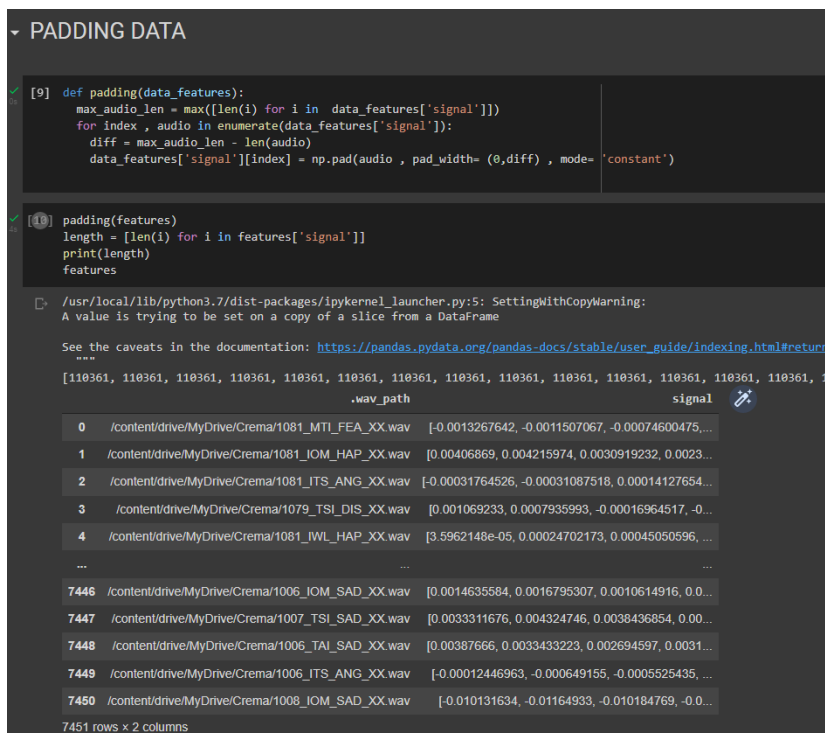
	.wav_path	signal	labels
0	/content/drive/MyDrive/Crema/1081_MTI_FEA_XX.wav	[-0.0013267642, -0.0011507067, -0.00074600475, ...	FEA
1	/content/drive/MyDrive/Crema/1081_IOM_HAP_XX.wav	[0.00406869, 0.004215974, 0.0030919232, 0.0023...	HAP
2	/content/drive/MyDrive/Crema/1081_ITS_ANG_XX.wav	[-0.00031764526, -0.00031087518, 0.00014127654...	ANG
3	/content/drive/MyDrive/Crema/1079_TSI_DIS_XX.wav	[0.001069233, 0.0007935993, -0.00016964517, -0...	DIS
4	/content/drive/MyDrive/Crema/1081_IWL_HAP_XX.wav	[3.5962148e-05, 0.00024702173, 0.00045050596, ...	HAP
...
7446	/content/drive/MyDrive/Crema/1006_IOM_SAD_XX.wav	[0.0014635584, 0.0016795307, 0.0010614916, 0.0...	SAD
7447	/content/drive/MyDrive/Crema/1007_TSI_SAD_XX.wav	[0.0033311676, 0.004324746, 0.0038436854, 0.00...	SAD
7448	/content/drive/MyDrive/Crema/1006_TAI_SAD_XX.wav	[0.00387666, 0.0033433223, 0.002694597, 0.0031...	SAD
7449	/content/drive/MyDrive/Crema/1006_ITS_ANG_XX.wav	[-0.00012446963, -0.000649155, -0.0005525435, ...	ANG
7450	/content/drive/MyDrive/Crema/1008_IOM_SAD_XX.wav	[-0.010131634, -0.01164933, -0.010184769, -0.0...	SAD

7451 rows x 3 columns

Play and plot the audios:



Some work on data



2-Creating the feature spaces:

```
def create_feature_space(audio_data_frame, sampling_rate):
    zero_crossing_rate_Time = []
    signal_Energy = []
    mel_spectrogram = []

    for audio in audio_data_frame['signal']:
        zero_crossing_rate_Time.append(lib.feature.zero_crossing_rate(y = audio))
        signal_Energy.append(lib.feature.rms(y = audio))
        mel_spectrogram.append(lib.feature.melspectrogram(y= audio, sr= sampling_rate, n_mels= 64))

    zero_crossing_rate_Time = pd.Series(zero_crossing_rate_Time, name= 'zero_crossing_rate_Time')
    signal_Energy = pd.Series(signal_Energy, name= 'Energy')
    mel_spectrogram = pd.Series(mel_spectrogram, name= 'mel_spectrogram')
    _audio_data_frame = pd.concat([_audio_data_frame, zero_crossing_rate_Time, signal_Energy, mel_spectrogram], axis=1)

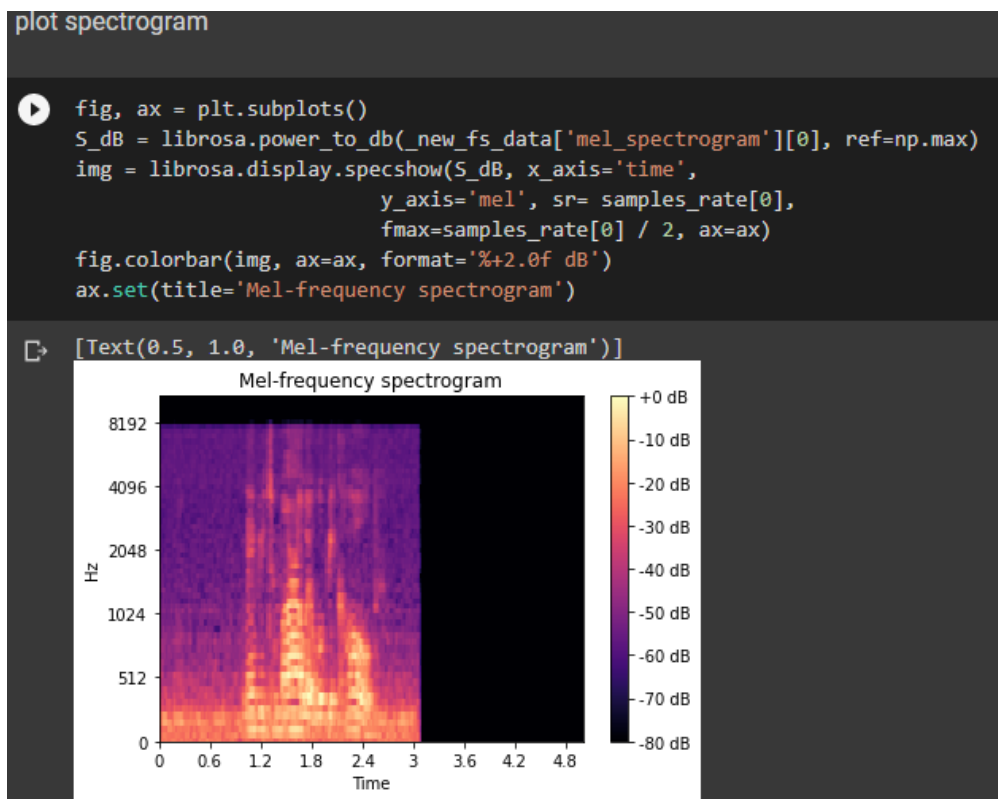
    return _audio_data_frame

[70] _new_fs_data.to_pickle('/content/drive/MyDrive/feature_space.pkl')
```

```
if not os.path.exists(r'/content/feature_space.pkl'):
    _new_fs_data = create_feature_space(features, samples_rate[0])
    _new_fs_data.to_pickle('/content/feature_space.pkl')
else:
    _new_fs_data = pd.read_pickle('/content/feature_space.pkl')
```

	.wav_path	signal	zero_crossing_rate_Time	Energy	mel_spectrogram
0	/content/drive/MyDrive/Crema/1081_MTI_FEA_XX.wav	[-0.0013267642, -0.0011507067, -0.00074600475, ...]	[[0.01171875, 0.0166015625, 0.0234375, 0.03076, ...]]	[[0.003751567, 0.00499682, 0.005775987, 0.0059, ...]]	[[0.041217044, 0.036117308, 0.041154295, 0.050, ...]]
1	/content/drive/MyDrive/Crema/1081_IOM_HAP_XX.wav	[0.00406869, 0.004215974, 0.0030919232, 0.0023, ...]	[[0.01220703125, 0.021484375, 0.0322265625, 0, ...]]	[[0.005943279, 0.0060860105, 0.0055265073, 0, ...]]	[[0.045877956, 0.018270325, 0.03386492, 0.0553, ...]]
2	/content/drive/MyDrive/Crema/1081_ITS_ANG_XX.wav	[-0.00031764526, -0.00031087518, 0.00014127654, ...]	[[0.009765625, 0.0224609375, 0.03076171875, 0, ...]]	[[0.006384029, 0.0061387853, 0.005591352, 0, ...]]	[[0.07248108, 0.064830616, 0.04831033, 0.06275, ...]]
3	/content/drive/MyDrive/Crema/1079_TSI_DIS_XX.wav	[0.001069223, 0.0007935993, -0.00016964517, 0, ...]	[[0.0166015625, 0.02294921875, 0.02783203125, ...]]	[[0.0044570197, 0.0051914803, 0.006889936, 0, ...]]	[[0.07827319, 0.029487826, 0.035867628, 0.0555, ...]]
4	/content/drive/MyDrive/Crema/1081_IWL_HAP_XX.wav	[3.5962148e-05, 0.00024702173, 0.00045050596, ...]	[[0.02294921875, 0.03466796875, 0.03955078125, ...]]	[[0.0047711073, 0.004586294, 0.0052237143, 0, ...]]	[[0.041727513, 0.07174656, 0.05327322, 0.06390, ...]]
...
7446	/content/drive/MyDrive/Crema/1006_IOM_SAD_XX.wav	[0.0014635584, 0.0016795307, 0.0010614916, 0, ...]	[[0.01513671875, 0.02685546875, 0.04541015625, ...]]	[[0.005884879, 0.00562598, 0.004991079, 0.0050, ...]]	[[0.018325038, 0.04317272, 0.05460701, 0.04703, ...]]
7447	/content/drive/MyDrive/Crema/1007_TSI_SAD_XX.wav	[0.0033311676, 0.004324746, 0.0038436854, 0, ...]	[[0.01220703125, 0.01708984375, 0.0244140625, ...]]	[[0.0071519525, 0.007694964, 0.0076097036, 0, ...]]	[[0.006080256, 0.010732218, 0.014269074, 0.025, ...]]

Mel spectrogram example:



3-Building models 1D-CNN , 2D-CNN:

- Data split 30% test and 70% train
- 5% of train validation test

3) DATA SPLITTING

70% trainig set and 30% testing set 5% of trainig set will be validation set

```
[39] from sklearn.model_selection import train_test_split

freq_features = pd.DataFrame(_new_fs_data[['mel_spectrogram']])

#SPLIT DATA INTO 30% TEST AND 70% TRAIN
train_data_time , test_data_time , train_labels_time , test_labels_time = train_test_split(time_feature_space , labels , test_size= 0.3 , stratify= labels)
train_data_freq , test_data_freq , train_labels_freq , test_labels_freq = train_test_split(freq_features , labels , test_size= 0.3 , stratify= labels)

#SPLIT TRAIN INTO 5% VALIDATION AND 95% TRAIN
training_data_time , validation_data_time , trainig_labels_time , validation_labels_time = train_test_split(train_data_time , train_labels_time , test_size= 0.05 , stratify= train_labels_time)
training_data_freq , validation_data_freq , trainig_labels_freq , validation_labels_freq = train_test_split(train_data_freq , train_labels_freq , test_size= 0.05 , stratify= train_labels_freq)

training_data_time.shape , test_data_time.shape , validation_data_time.shape , freq_features.shape

((4954, 432), (2236, 432), (261, 432), (7451, 1))
```

1D – Model:

```
#WORK ON DATA IN TIME DOMAIN
#REQ: TO SEND THE DATA IN TIME DOMIAN
with tf.device(tf.test.gpu_device_name()):
    model_time = Sequential()
    model_time.add(Conv1D(512, kernel_size=5, padding='same', strides=1, activation='relu', input_shape=(x_trainig_time.shape[1],1)))
    model_time.add(BatchNormalization())
    model_time.add(MaxPooling1D(pool_size=5, padding='same', strides=2))

    model_time.add(Conv1D(512, kernel_size=5, padding='same', strides=1))
    model_time.add(BatchNormalization())
    model_time.add(MaxPooling1D(pool_size=5, padding='same', strides=2))

    model_time.add(Conv1D(128, kernel_size=5, padding='same', strides=1))
    model_time.add(BatchNormalization())
    model_time.add(MaxPooling1D(pool_size = 5, padding='same', strides=2))

    model_time.add(Flatten())
    model_time.add(Dense(256, activation='relu'))
    model_time.add(Dropout(0.3))
    model_time.add(Dense(6, activation='softmax'))
    model_time.compile(optimizer="rmsprop", loss=tf.keras.losses.CategoricalCrossentropy(), metrics=["accuracy", f1_m])

    model_time.summary()

[ ] model_time = tf.keras.models.load_model('/content/model_checkpoints/time_model_60')

[ ] history = model_time.fit(x_trainig_time, trainig_labels_time, validation_data=(x_val_time, validation_labels_time),
                             epochs=EPOCHS, batch_size=32,
                             callbacks=[earlystopping, learning_rate_reduction, cp_callback])
```

```
[ ] from sklearn.metrics import confusion_matrix , ConfusionMatrixDisplay


test_acc = model_time.evaluate(x_test_time , test_labels_time)[1]
print("Test Accuracy = ", test_acc * 100 , "%")

70/70 [=====] - 26s 364ms/step - loss: 1.4206 - accuracy: 0.4915 - f1_m: 0.4181
Test Accuracy = 49.150267243385315 %

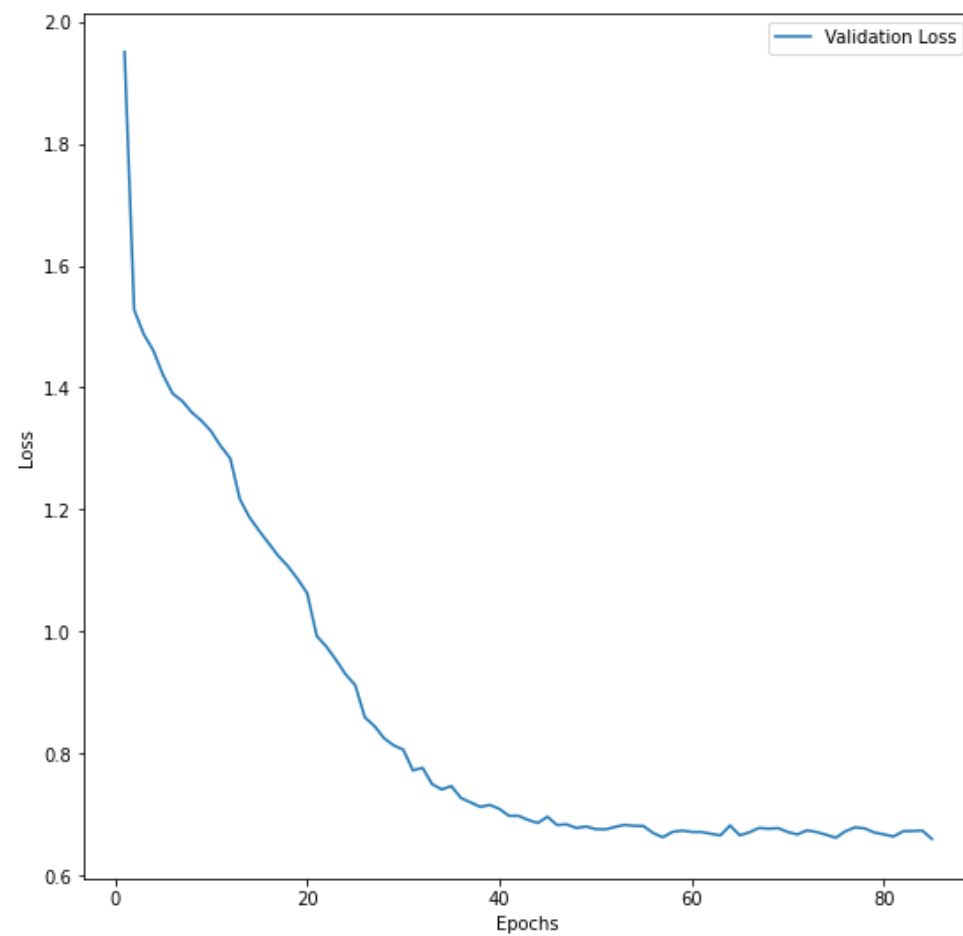

[ ] test_prediction = np.argmax(model_time.predict(x_test_time) , axis=1)
test_true = np.argmax(test_labels_time , axis=1)
conf_matrix = confusion_matrix(y_pred= test_prediction , y_true= test_true)

[ ] import seaborn as snNew
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = [9, 9]

snNew.heatmap(conf_matrix, annot=True , xticklabels= encode.classes_ , yticklabels= encode.classes_,fmt="d")
plt.show()
```



Most confusing class is FEA



2D-Model:

LOAD MODEL 2D FROM DRIVE

```
[46] model_spect= tf.keras.models.load_model('/content/checkpoints/spec')

[61] train_data_freq = training_data_freq['mel_spectrogram'].reset_index(drop=True)
      train_data_freq = train_data_freq.to_list()
      train_data_freq = tf.constant(train_data_freq)

[62] val_data_freq = validation_data_freq['mel_spectrogram'].reset_index(drop=True)
      val_data_freq = val_data_freq.to_list()
      val_data_freq = tf.constant(val_data_freq)

[63] tst_data_freq = test_data_freq['mel_spectrogram'].reset_index(drop=True)
      tst_data_freq = tst_data_freq.to_list()
      tst_data_freq = tf.constant(tst_data_freq)
```

```
[ ] train_data_freq.shape

TensorShape([4954, 64, 216])
```

```
[ ] del data
    del time_feature_space
```

REMOVE UNNECESSARY VARIABLES FROM RAM

```
[77] import gc
      with tf.device('/device:GPU:0'):
          for _ in range(100):
              gc.collect()
              tf.keras.backend.clear_session()

[82] with tf.device('/device:GPU:0'):
      model_spect.fit(train_data_freq, train_labels_freq,
                      validation_data=(val_data_freq, validation_labels_freq),
                      epochs=EPOCHS,
                      batch_size=BATCH_SIZE,
                      callbacks=[earlystopping, learning_rate_reduction, cp_callback])
```

```
[144] from sklearn.metrics import confusion_matrix , ConfusionMatrixDisplay
      with tf.device('/device:GPU:0'):
          test_acc_2D = model_spect.evaluate(tst_data_freq , test_labels_freq)[1]
          print("Test Accuracy = ", test_acc_2D * 100 , "%")


70/70 [=====] - 8s 108ms/step - loss: 1.4375 - accuracy: 0.5371 - f1_m: 0.5082
Test Accuracy = 53.71198654174805 %
```

```
[57] test_labels_freq

array([[0., 0., 0., 1., 0., 0.],
       [0., 0., 0., 0., 1., 0.],
       [0., 0., 0., 0., 0., 1.],
       ...,
       [0., 0., 1., 0., 0., 0.],
       [0., 0., 0., 0., 1., 0.],
       [0., 0., 0., 0., 1., 0.]], dtype=float32)
```




```
loss_values = history['loss']
epochs = range(1, len(loss_values)+1)
plt.plot(epochs, loss_values, label='Training Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```




```
[135] valid_loss = history['val_loss']
plt.plot(epochs, loss_values, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()
```



```
import seaborn as snNew
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = [9, 9]

snNew.heatmap(conf_matrix_freq, annot=True , xticklabels= encode.classes_ , yticklabels= encode.classes_,fmt="d")
plt.show()
```



Most confusing class HAP